

Meta Musical Memes

Megan Jasek, James King, Sean Underwood

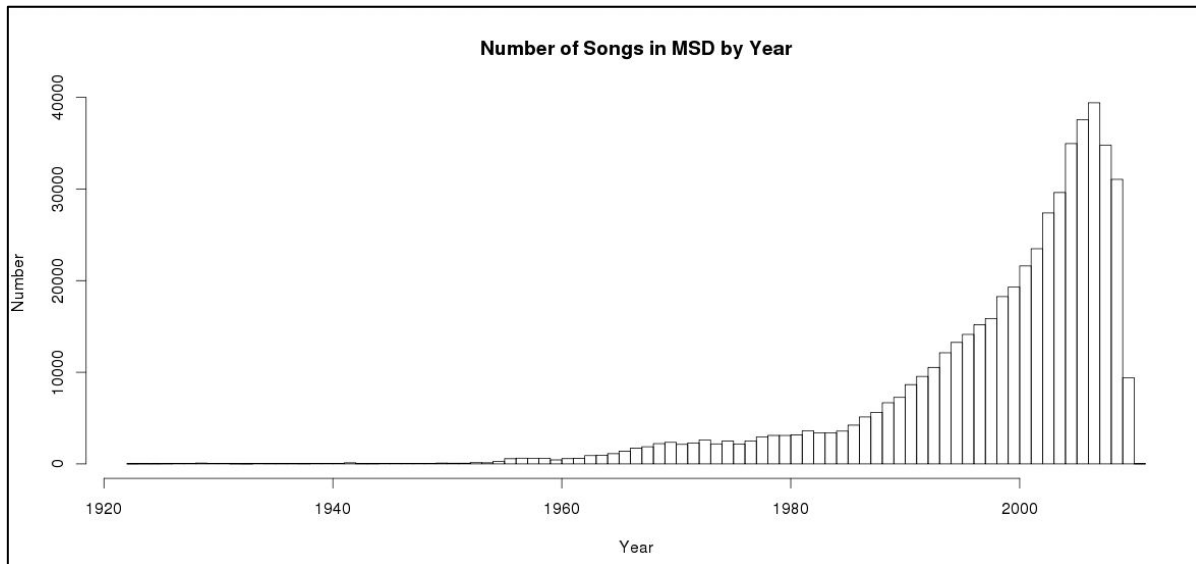
W205 Data Storage and Retrieval - Fall 2015

Problem Definition

Is there a basic formula that helps predict a song's commercial success?

Data

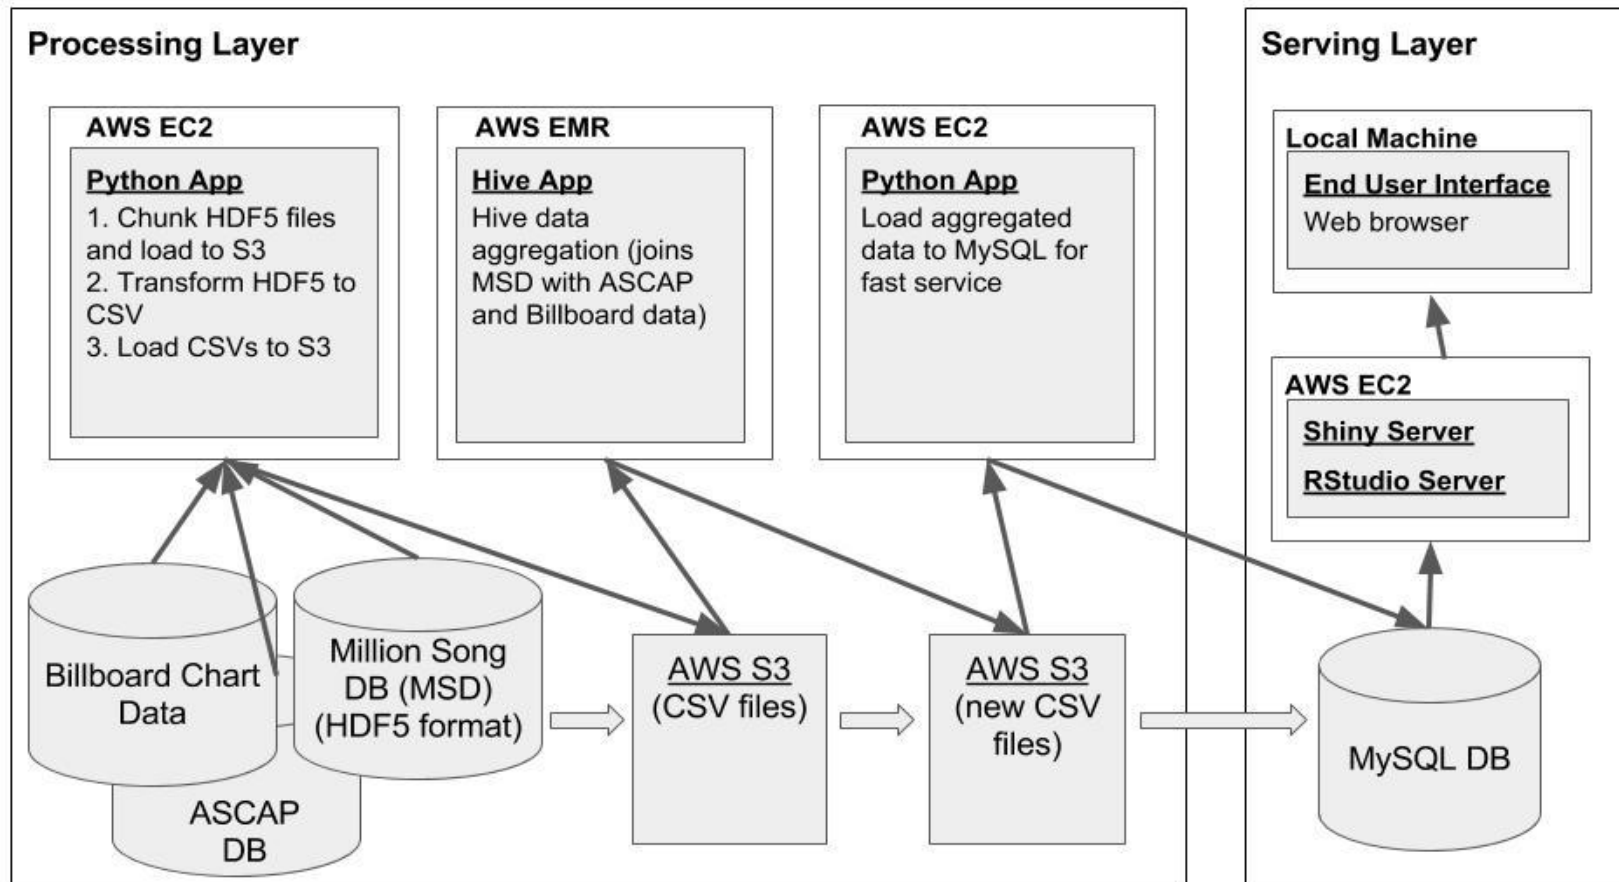
- **Million Song Dataset**
 - Song credits
 - Music theory info (key, tempo, etc.)
 - Audio features
 - Popularity estimates (“Hottnesss”)
- **Billboard charts**
 - Performer credits
 - Peak position
 - Peak year
- **ASCAP records**
 - Writer credits



Data - Acquisition and Organization

- Million Song Dataset:
 - Provided by Columbia University's LabROSA
 - Publicly available on Amazon AWS as an EBS Snapshot
 - File format: HDF5 (one file per song)
 - 500 GB
- Billboard Data
 - Fan-maintained
 - Excel format
 - 29 MB
- ASCAP Data
 - Provided by ASCAP
 - CSV format
 - 1.6 GB

Architecture



Processing Layer

- ETL: Million Song Dataset (500 GB)
 - HDF5 - Does not interact well with HDFS.
 - HDF5 is optimized for high performance when being used from a single process that does not need to load the whole file into memory at once.
 - When a file gets split across a block boundary, there's no good way for workers to re-assemble the data.
 - h5py python library
 - Extraction - multiprocessing python library - run code in parallel
 - xLarge AWS instance with 8 VCPUs
 - Still took 12 hours to extract data
- ETL: Billboard Data (30 MB)
 - MS Excel file
 - Extracted 4 attributes (song, artist, peak position, peak year) using python
- ETL: ASCAP Data (1.5 GB)
 - CSV file in long format, unusual name formatting: "last middle first"
 - Extracted 3 attributes (song, artist, writer)

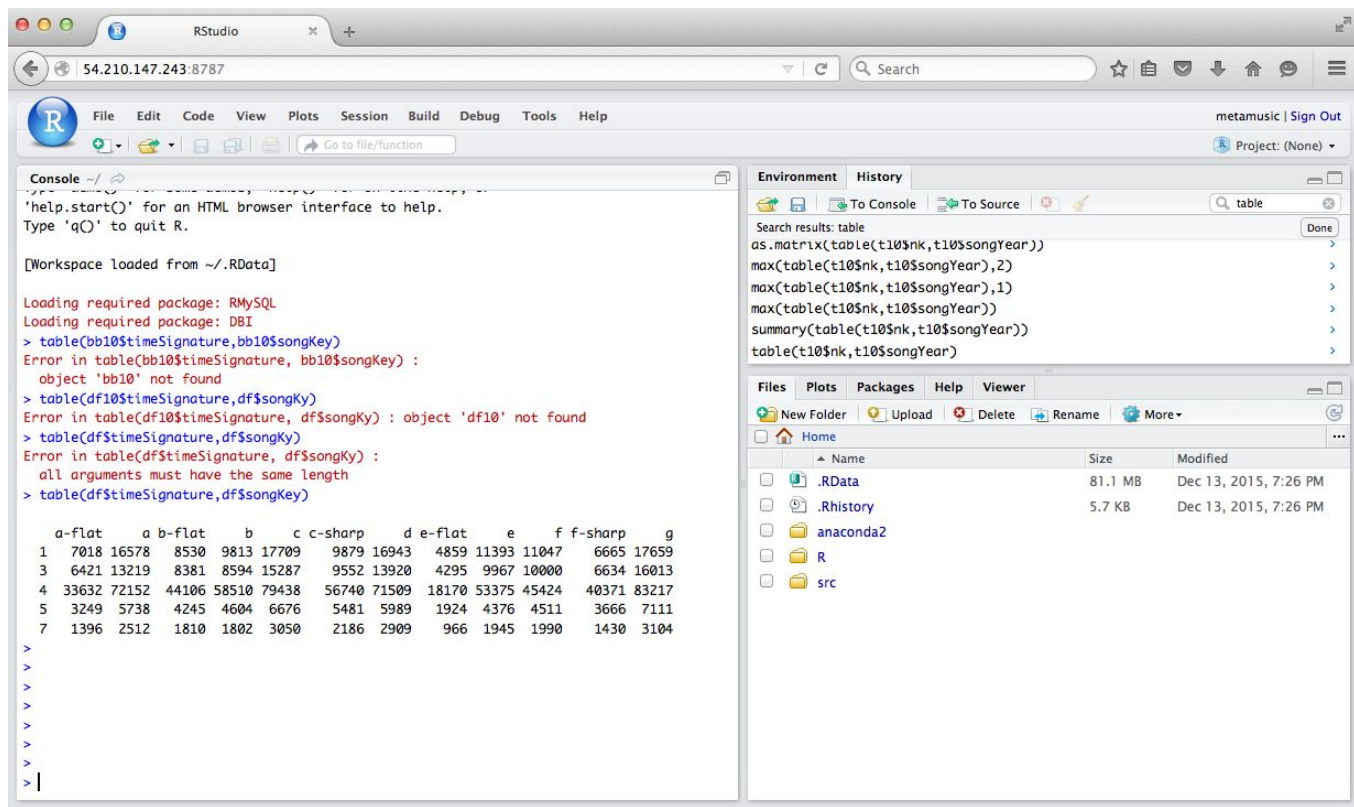
Processing Layer (cont.)

- Data Merge
 - All data converted to CSV format
 - Uploaded to Amazon S3
 - Amazon Elastic Map Reduce (EMR)
 - Hive job to clean and merge datasets
 - Massaged MSD 0 values
 - Converted numeric encodings for key and mode to text
 - Output uploaded to Amazon S3
 - 3 m3.xlarge instances - total time 5 minutes
 - Hive trickiness
 - Hive does not proactively validate data
 - It's very important to manually validate
- ASCAP: There's no escape from incomplete data
 - We failed to find a reliable way to bridge the lack of performer credits in the ASCAP data.
- Produced 2 datasets
 - One with ASCAP and one without

Serving Layer

- Merged data in MySQL DB
- Interactive Output
 - RStudio Server
 - Running on EC2 instance
 - Exposed full R ecosystem in a web browser
 - Shiny
 - Toolkit for writing web applications in R
 - Create interactive UI widgets hooked up to R commands
 - Shiny Server
 - Running on EC2 instance
 - Exposes Shiny applications in a web browser

RStudio Server



The screenshot displays the RStudio Server web interface. The top navigation bar includes menus for File, Edit, Code, View, Plots, Session, Build, Debug, Tools, and Help. The address bar shows the URL 54.210.147.243:8787. The console on the left shows the following output:

```
'help.start()' for an HTML browser interface to help.
Type 'q()' to quit R.

[Workspace loaded from ~/.RData]

Loading required package: RMySQL
Loading required package: DBI
> table(bb10$timeSignature, bb10$songKey)
Error in table(bb10$timeSignature, bb10$songKey) :
  object 'bb10' not found
> table(df10$timeSignature, df10$songKey)
Error in table(df10$timeSignature, df10$songKey) : object 'df10' not found
> table(df$timeSignature, df$songKey)
Error in table(df$timeSignature, df$songKey) :
  all arguments must have the same length
> table(df$timeSignature, df$songKey)
```

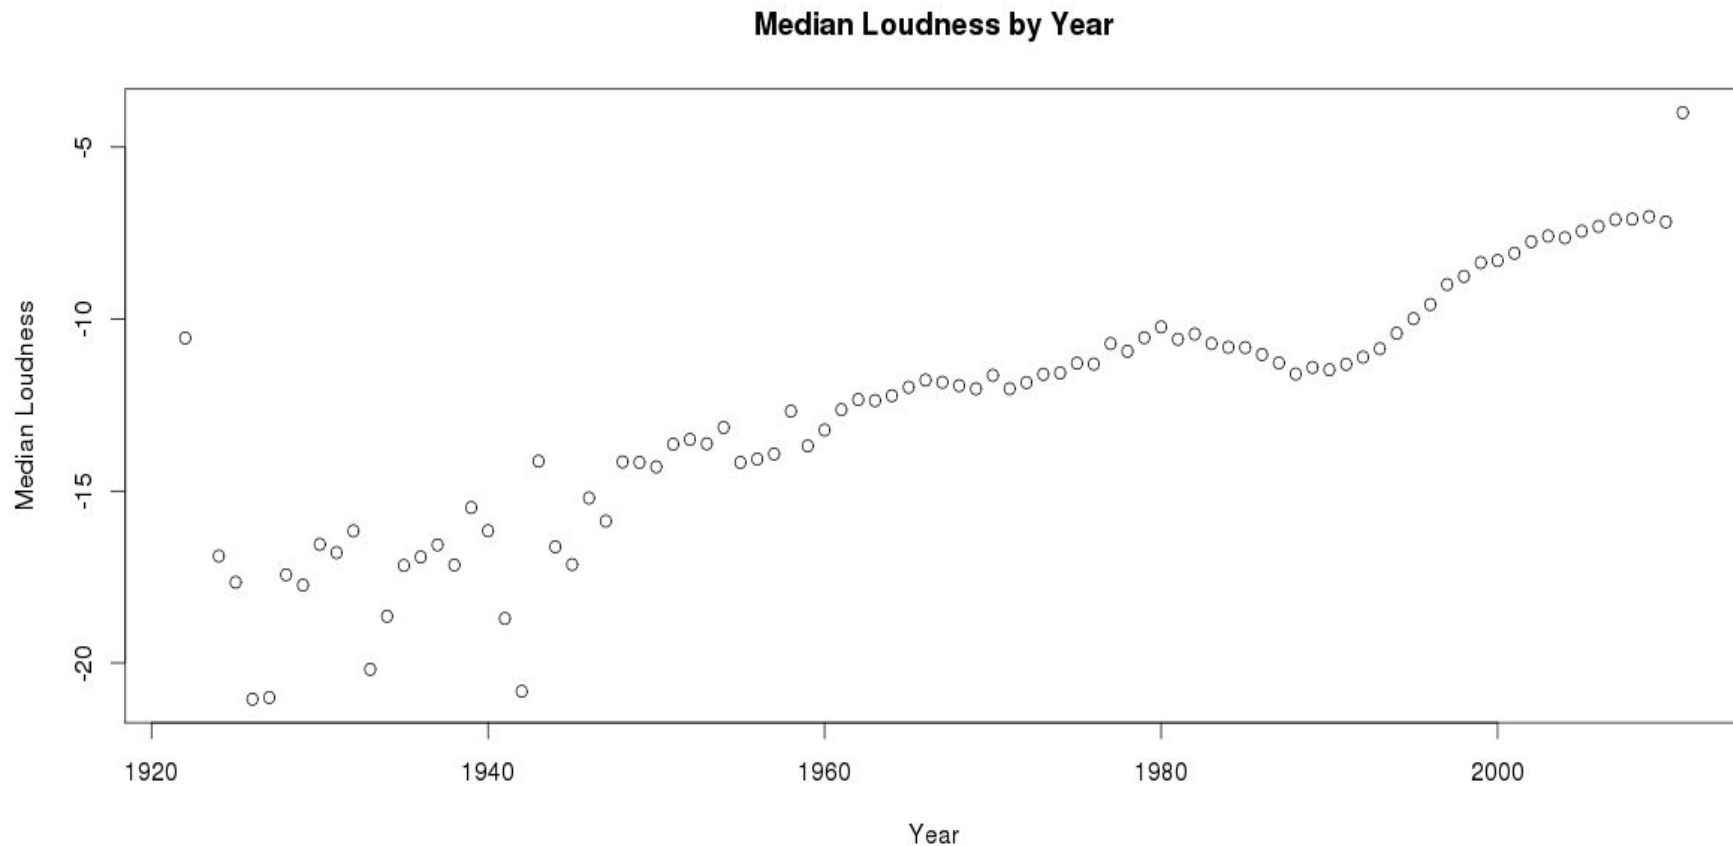
The right-hand pane shows the Environment and History tabs. The Environment tab displays search results for the variable 'table':

```
Search results: table
as.matrix(table(t10$nk, t10$songYear))
max(table(t10$nk, t10$songYear), 2)
max(table(t10$nk, t10$songYear), 1)
max(table(t10$nk, t10$songYear))
summary(table(t10$nk, t10$songYear))
table(t10$nk, t10$songYear)
```

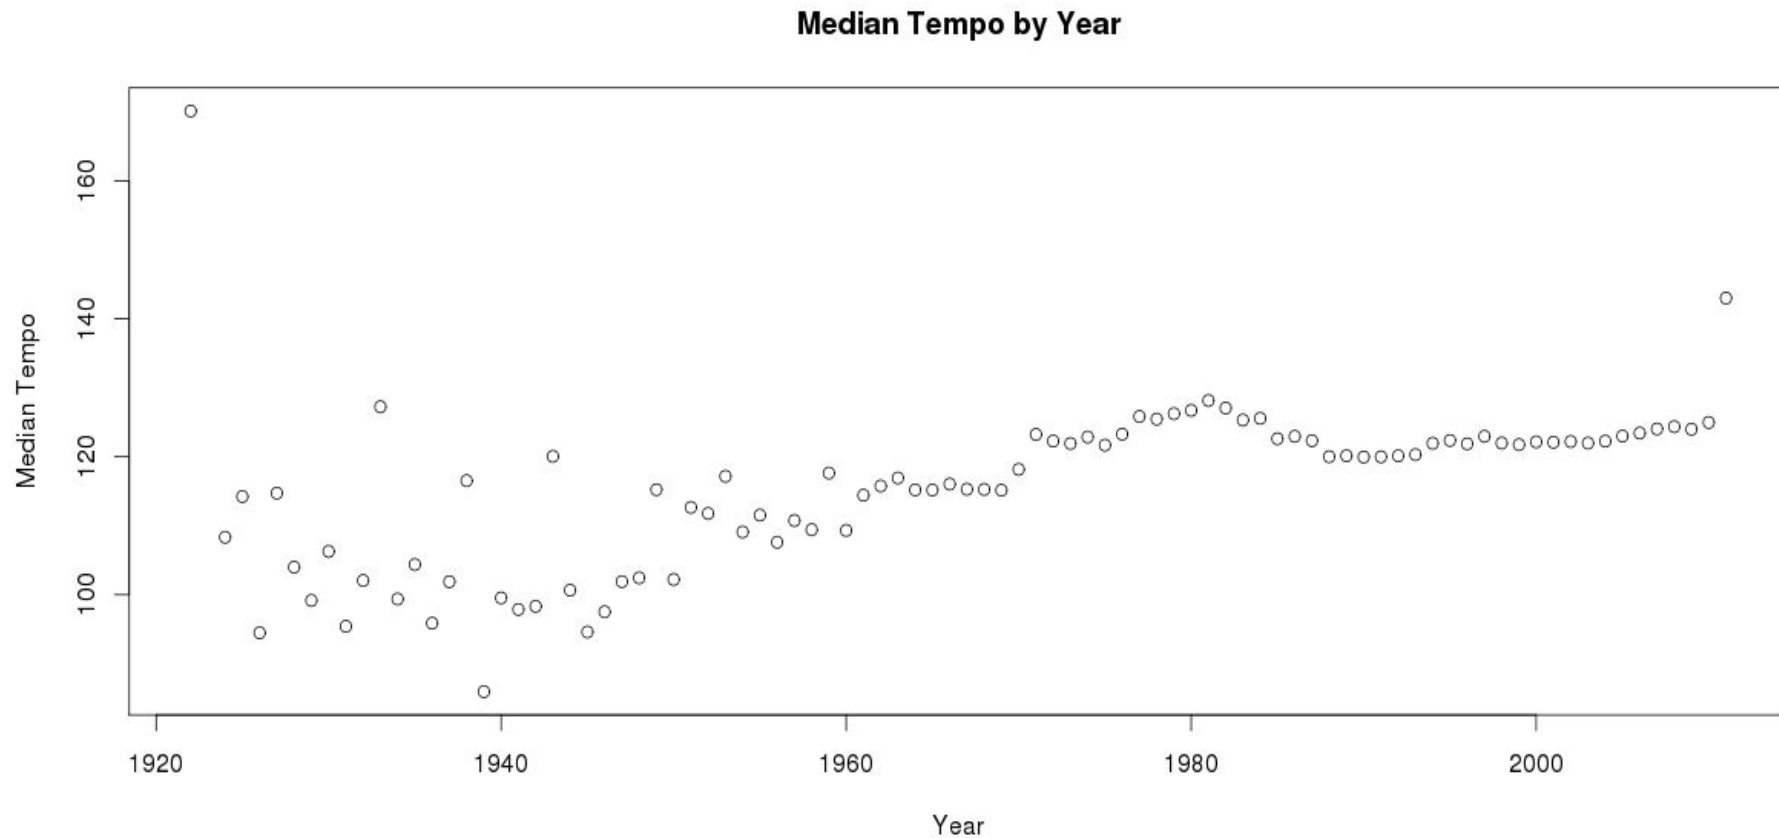
Below the Environment tab is a file explorer showing the file structure of the project:

Name	Size	Modified
.RData	81.1 MB	Dec 13, 2015, 7:26 PM
.Rhistory	5.7 KB	Dec 13, 2015, 7:26 PM
anaconda2		
R		
src		

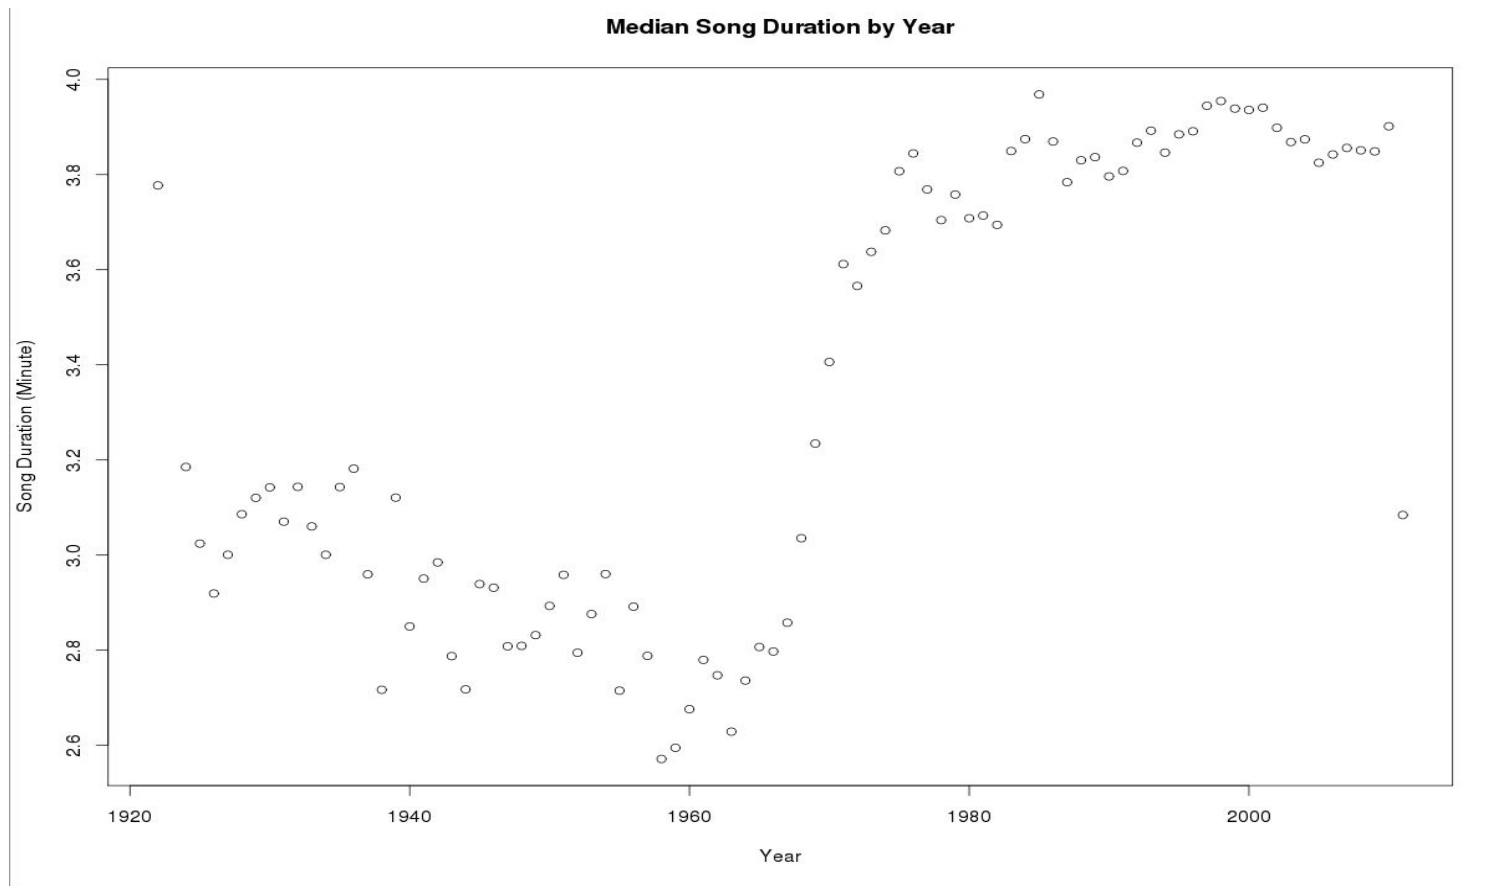
RStudio Results: Songs are Getting Louder.....What?



RStudio Results: Songs. Aren't. Getting. Any. Faster.




RStudio Results: Songs are Getting Lonnngnnnnnger...



Analysis: Shiny

Example 1:



Song Statistics by Artist



Analysis: Shiny

Example 2:

Song Hotness Regression

 52.91.231.109:3838/song_hotness_regression/

Choose the Variables for Regression for Song Hottness
☒ artistHottness
☐ duration
☐ tempo
☐ songKey
☐ loudness
☐ timeSignature
☒ peakPosition
☐ songYear
☐ billboardYear
☐ mode

Call:
lm(formula = as.formula(paste("songHottness ~", terms)), data = df)

Residuals:

Min	1Q	Median	3Q	Max
-0.69133	-0.10586	0.02192	0.12318	0.46395


Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	4.136e-01	6.876e-03	60.16	<2e-16 ***
artistHottness	5.344e-01	1.311e-02	40.77	<2e-16 ***
peakPosition	-9.584e-04	5.599e-05	-17.12	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.1683 on 10271 degrees of freedom
(1990344 observations deleted due to missingness)
Multiple R-squared: 0.1625, Adjusted R-squared: 0.1624
F-statistic: 996.7 on 2 and 10271 DF, p-value: < 2.2e-16

Analysis: Shiny Example 3: Scatter Plots



Architecture - Alternatives

- Processing Layer

- Leave the single-song HDF5 files as-is
 - The “small files problem” wasn’t nearly as problematic as the Hadoop/HDF5 problem.
- Run the entire job in Spark
- Convert HDF5 files to another format (e.g., JSON) from the outset
 - Human-readable
 - Hadoopable
 - Bigger data size, but ㄟ(っ)ㄟ

- Serving Layer

- DB - any relational database
- BlinkDB - sampling database to speed up R visualizations
- Tableau
- Python with visualizations in D3

Future of Project - Scaling

- As expected, the Hive portions of the job scale very well.
- HDF5 file processing was largely sequential.
 - Scaling would require devising a way to store HDF5 files in a data lake
 - Alternatively, they could be converted to a more compatible format at import time
- Shiny had difficulty with the volume of data
 - R loads all data points into RAM, which imposes scaling limits
 - Regression and plotting of such large volumes of data is slow
 - Could instead use a sampling database such as BlinkDB to control the amount of data